

# Heuristics

Matthew Guzdial

# A\*

add **start** to **openSet**

while **openSet** is not empty:

*current* = **openSet**.pop()

    if *current* == **goal**:

        return reconstruct\_path(*current*)

**closedSet**.Add(*current*)

    for each *neighbor* of *current*:

        if *neighbor* in **closedSet**:

            continue

*gScore* = *current.gScore* + heuristic(*current*, *neighbor*)

        if *neighbor* not in **openSet**:

**openSet**.add(*neighbor*)

        else if *gScore* < **openSet**.get(*neighbor*).*gScore*

**openSet**.replace(**openSet**.get(*neighbor*), *neighbor*)

# Heuristic


- A function that given a state and goal, estimates remaining distance
- $A^*$  is guaranteed to find the optimal path to the goal *if* the heuristic ( $h$ ) is admissible

# What does admissible mean?

Some heuristic ( $h$ ) is admissible iff for all states  $n$ ,  $h(n) \leq h^*(n)$


- Where  $h^*$  is the true distance

# Example: Navigating Roads



- $h^*(B) = 2$
- $h^*(C) = 1.5$
  
- $h(B) = 2$
- $h(C) = 24$

# Alternative Heuristic



- $h^*(B) = 2$
- $h^*(C) = 1.5$

If  $h$ =Straight line distance

- $h(B) = 1.5$
- $h(C) = 1.5$
- Roads can never be shorter than this

# You must Prove Admissibility

- If your heuristic is not admissible, we can't guarantee that  $A^*$  gives an optimal solution
- Why?
- Can you think of a case where we *wouldn't* care if our  $A^*$  solution was non-optimal?

# 8-Puzzle Game



☐ Show numbers

Shuffle

Solve



## Example 2: 8 Puzzle

7	2	4
5		6
8	3	1



	1	2
3	4	5
6	7	8

What would be a good heuristic?

# Options

## 1. #1 tiles out of place

- Need at least this many moves
- Best case: each tile is exactly 1 out of place
- ...But probably need many other moves

## 2. Manhattan Distance

- # of square blocks you have to walk on a grid
- (Same justification as above)

# How do we choose between heuristics?

Given last slide's two heuristics  $h_1$  and  $h_2$

For initial state  $s_1$

$$h_1(s_1) = 8$$

$$h_2(s_1) = 18$$

$$h^*(s_1) = 26$$

$h_2$  better than  $h_1$

Why?

$h_2$  is closer to  $h^*$  but for all states  $n$ ,  $h_2(n) \leq h^*(n)$

# Informedness

- Degree to which a heuristic  $h$  captures insight
- Informedness = size of total search space / avg. # of states explored with heuristic  $h$
- Do we want high or low informedness?

# Dominance

$h_a$  dominates  $h_b$  when  $h_b(n) \leq h_a(n) \leq h^*(n)$

- $h_b$  underestimates more
- $h_a$  and  $h_b$  are both admissible

# Can you think of a heuristic that is always admissible?

- Yes  $h(n) = 0$  for all states  $n$
- But not *informed*
- Admissibility does not guarantee a heuristic is any good (only optimality of solution)
- A non-admissible heuristic can be more informed (but not dominate or admissible), which will give better average time performance

# Consistency

- AKA monotonicity: Always increasing/decreasing
- Admissibility: never overestimate
- Overestimates cause shortcuts



# What happens if we underestimate too much?

- Still admissible
- May be fooled into spending too much time in useless part of space
- Will eventually find optimal solution though
- $h(n) = 0$  is just breadth first search

# Okay cool, but how do I come up with a heuristic for a problem?

1. In cases where runtime is a concern, should be  $O(\text{Polynomial})$  or better
  - Relax the problem
    - At any given state, make a simpler problem
    - Solve simpler problem exactly
    - Return true cost of simpler problem

# Example: 8 Puzzle

What are the constraints?

Tile can move from square A to square B if:

1. A adjacent to B
2. B is blank

7	2	4
5		6
8	3	1

# What if $h == h^*$ ?

Is it still admissible?

- Yes of course!

So we should always strive for that right?

- Nope, if  $h$  is too expensive

# Heuristics Review

- Admissibility: all states  $n$ ,  $h(n) \leq h^*(n)$ 
  - Must be true to guarantee optimal solution
- Informedness: measure of how little we need to explore (total # of states/# of states we explore)
- Dominance:  $h_2$  dominates  $h_1$ ,  $h_1(n) \leq h_2(n) \leq h^*(n)$
- Consistency: Always increasing or decreasing relative to neighbors

Search

# Randomized Optimization

- Intuition: What if we have a heuristic but no goal, how do we pick the best final state?
- Each state is now some representation
  - Note: I say “representations” here, but we can think about literally any data structure here
  - Almost any knowledge representation can be searched over, given some definition of “best”.

# Naïve Approach: Generate and Test

Create representation at random and then see how well it performs.

Issues:


- Large hypothesis spaces
- Doesn't learn anything



# Randomized Optimization Approaches

- Hill-climbing aka Greedy Search
- Simulated Annealing
- Genetic Algorithms

# Hill Climbing




If y-axis is some heuristic and x-axis are representations, we can imagine a pretty simple algorithm to find best representation

# Hill Climbing Pseudocode


```
curr = random_selection_from_space()  
score = heuristic(curr)  
prevScore = 0  
while score > prevScore:  
    prevScore = score  
    for n in neighbors(curr):  
        if heuristic(curr) < heuristic (n):  
            curr = n  
            score = heuristic(curr)  
return curr
```

# Example


1.) Pick a random starting location




2.) Find Neighbors




3.) Select best neighbor or finish



Repeat 2 and 3 till....



# Hill Climbing Exercise



Given input points 1-5, what points A-E would they end up at following the hill climbing algorithm?

# Answer

1. A
2. A
3. C
4. D
5. ??? (depends on neighbor function)
  - e.g. Left or right first? one step or multiple?

# Local vs. Global Maxima

- Local Maximum: A point with only worse neighbors
- Global Maximum: A point better than all other points
  - What we really want

# Issues

- Choice of neighbor function
  - Deciding between ties
  - What neighbors do we consider
- Local vs. Global Maxima
  - Can we ensure we get a global maxima?
  - What are ways we can improve our chances?



# Choice of Neighbor Function

- It's easy to “see” where to go next if we visualize the entire space.
- But functionally we do not have the entire space already, we generate each neighbor as needed
- Show example on board if they don't get it.

# Neighbor Function Exercise

**Representation:** 6 bit strings (ex. 101010)

**Heuristic:** Number of shared bits with 101010

**Neighbor Function 1:** Flip pairs of adjacent bits

100000 → 111000 or 010000 but not 101010

**Neighbor Function 2:** Flip any two bits

100000 → 111000 or 010000 or 101010

Q: What's the maxima we'll reach with the random starting points 010101, 000100, and 110000?

# Answers

Q: What's the maxima we'll reach with the random starting points 010101, 000100, and 110000?



N1: 101010 (6), 000010 (4), and 101000 (5)

N2: 101010 (6), 101010 (6), and 111010/101000(5)

# Finding the Global Maximum

- One simple idea: Hill-climbing with random restarts
- Keep running hill-climbing for a given number of times or we stop seeing improvements, starting at a random place each time

# Still not quite what we want



Random restarts not likely to work well in some spaces (hilly or skewed spaces)

# Bad Solutions


What if we could better check the entirety of the available space of representations?

- **Exhaustive Search:** Check every single point, grab the best (infinite/large  $H$  problem again)
- **Random Walk:** Keep making random moves till we stopping seeing anything better (sometimes better than exhaustive)

# Simulated Annealing

Intuition: What if we sometimes took random steps at the start of search and took less and less as the search continued?

- We could better **explore** the space



# Simulated Annealing

Think of it as a hybrid between random walk and hill-climbing.

Inspired by annealing in metallurgy/  
blacksmithing





# Simulated Annealing Pseudocode

```
curr = random_selection_from_space()
score = heuristic(curr)
temperature = some_max_temperature
while temperature > 0 :
    next = random_neighbor(curr)
    temperature = schedule(temperature)
    changeH = heuristic(next) - score
    if (changeH > 0): curr = next, score = changeH + score
    else if ( $e^{(\text{changeH}/\text{temperature})} > \text{Rand}()$ ): curr = next, score = changeH + score
return curr
```

# Simulated Annealing Behavior

- If temperature =  $\infty$ , behavior is random walk
- If temperature = 0, behavior is hill-climbing
- Key is to choose a good schedule to guide decrease of temperature.

# Simulated Annealing Problem



Would Simulated Annealing do any better?  
Answer: Sometimes to the right, not the left


# Genetic Algorithms

# Genetic Algorithms in the Wild



2006 NASA Spacecraft Antenna

# Genetic Algorithms in the Wild




**DeepMind:** PathNet: Evolution Channels Gradient Descent in Super Neural Networks


# Genetic Algorithms

- Sometimes called evolutionary search
- Intuition: Instead of searching/walking from a single point -> many points spread across the search space
  - Use the knowledge gained from these points to take big, but likely better leaps

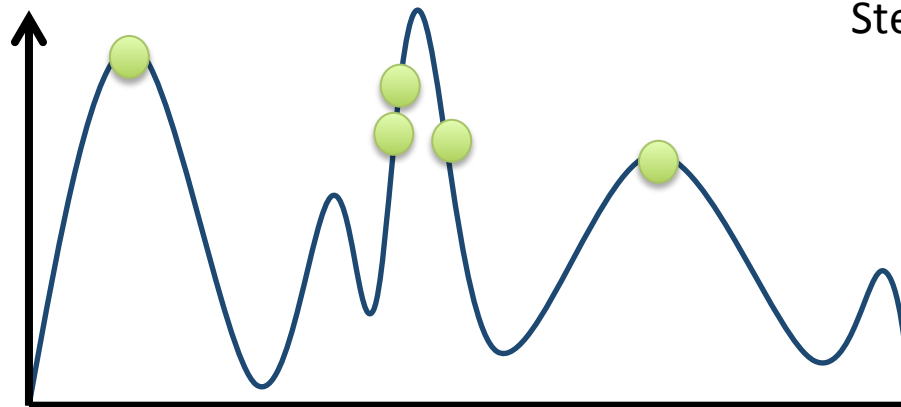
# Genetic Algorithms Visualized



Initialization



Step 2




Step 3



# Genetic Algorithms: How the heck?

- Inspired by biological evolution
- Population of points
  - Undergoes mutation (random walk)
  - Reproduces (crossover)
  - Fittest survive (heuristic)



# GA Pseudocode

population = set of random points of size  $n$

time = 0

while **heuristic**(population) < *threshold* and time < max:

    time ++

**Mutate**(population)

    population = **Crossover**(population)

    population = **Reduce**(population) *//back to size  $n$*

return  $\max_{\text{heuristic}}$ (population)

# Key Functions

- **Mutate:** This is a random walk, according to some user-set probability replace a member of the population with a random neighbor
- **Crossover:** Mix two members of population selected according to their heuristic value
  - Population  $\geq 2 \cdot n$
- **Reduce:** Take the  $n$  best, or take the  $n$  new best


# Crossover Strategies

Imagine we have the 6 bit strings from before. We could imagine many possible crossovers given two parents A and B.


- Single-point: Pick some index  $i$ , take all elements of A before  $i$ , all elements of B  $\geq i$ 
  - 100100 and 111111 with  $i=4 \rightarrow 100111$

# Strategies continued...


- Two-point: pick two indexes  $i, j$  where  $j > i$  and  $j < \text{length}$ . Take all points  $[0, i)$  from A, all points  $[i, j]$  from B and all points  $[j+1, l)$  from A
  - 100100 and 111111 with  $i=2, j=4 \rightarrow 101110$
- Uniform: For each index/variable, randomly select from both parents



Initialization



Step 2




Step 3

- What is  $n$ ?
- What is the likely crossover function?
- What is the reduce function?

# GA Exercise

**Problem:** plans for a grid world. Start (green), get to blue w/ minimum cost. Red is lava (infinite cost).

**Representation:** sequences of length 10, actions down (1 cost), right (1), left (1), up (1), and wait (0)



Sequence: down, right, right, right, right, right, right, right, right, down

- Would get to goal with 10 cost.

Give me population size and heuristic, mutation, crossover and reduce functions.

# GA Exercise (My Answers)

- Population size: 10 (arbitrary)
- Heuristic: Manhattan Distance to Goal – Cost
- Crossover: Single-point (first half of plan/  
second half of another plan)
- Reduce: Take 10 best of parents and children  
(don't want to lose good plans)



# Random Optimization Summary

- Hill-climbing: Strong simple approach, but weak to local maxima
- Simulated Annealing: With temperature we can get benefits of random walk and hill-climbing, but also some of their downsides
- Genetic Algorithms: With quite a bit of authoring we can get a robust solution to optimization problems
  - Called the 2<sup>nd</sup> best solution to anything